

Replicated end-to-end latency scaling

